

Scalable Real-time Detection of AI-Generated Arabic Text: A Distributed Data Pipeline Approach

Abdulmajeed Alqhtani
Taibah University
Almadinah, Saudi Arabia
AlqhtaniAbdulmaajeed@gmail.com

Abstract

The rapid proliferation of Large Language Models (LLMs) has introduced significant challenges to academic integrity, particularly concerning the undetected use of AI-generated text. This project proposes a scalable, Big Data-driven machine learning system designed to distinguish between human-written and AI-generated Arabic academic abstracts. Utilizing a dataset of 16,895 balanced records, the system was built end-to-end using **Apache Spark** to handle large-scale data ingestion, processing, and real-time streaming.

The core methodology relied on advanced feature engineering that combined standard NLP techniques with custom linguistic analysis. A **TF-IDF** vectorizer (2,000 features) was utilized to capture the foundational lexical baseline of the texts. To capture the "stylometric footprint" of AI versus human writers, five custom structural features—including word length variance, characters per paragraph, command verb counts, average sentence length, and a formality score—were extracted using the **Camel Tools** library. To overcome the computational bottlenecks associated with large-scale Arabic text processing, the engineered features were persisted using distributed **Parquet** storage, optimizing the memory footprint and breaking the processing lineage.

Three classification algorithms were trained and evaluated: Naïve Bayes (Baseline), Gradient Boosted Trees (GBT), and **Random Forest (Champion Model)**. The Random Forest classifier demonstrated superior performance, achieving an accuracy exceeding **90%**. A deep dive into feature importance revealed a critical insight: AI models exhibit highly standardized sentence structures. Among the custom stylometric variables, **Average Sentence Length (f84)** proved to be the most decisive discriminator, contributing 52.4% of the relative predictive power. Furthermore, the Command Count feature (f61) consistently registered zero variance, perfectly reflecting the objective, declarative nature of academic writing, which the model successfully learned to navigate.

Finally, to demonstrate the system's deployment readiness, a real-time **Spark Structured Streaming** pipeline was developed. The streaming simulation successfully ingested and classified incoming text files on the fly, proving that the architecture is not only highly accurate but also fully capable of conducting live, real-time academic auditing.

1. Methodology and Data Engineering

1.1. Data Ingestion and Preprocessing

The foundation of this project relies on a robust data ingestion pipeline built with Apache Spark. The dataset, sourced from the Hugging Face hub, consists of Arabic academic abstracts. To prevent class imbalance biases during model training, the dataset was strictly balanced, resulting in an aggregated corpus of approximately 17,000 records (evenly distributed between human-authored and AI-generated texts). Preprocessing involved standard Arabic Natural Language Processing (NLP) techniques, including diacritic removal, character normalization, and the elimination of non-alphabetic symbols, ensuring a clean lexical baseline for subsequent transformations.

1.2. Feature Engineering Strategy

A dual-track feature engineering approach was adopted to capture both the vocabulary and the structural "fingerprint" of the texts:

- **Lexical Features (TF-IDF):** A Term Frequency-Inverse Document Frequency (TF-IDF) vectorizer was applied to extract the top 2,000 contextual terms. This provided the models with a quantifiable understanding of the most discriminative vocabulary used by AI versus human writers.
- **Stylometric Features (Custom UDFs):** To capture the syntactic structures and writing styles, five custom User-Defined Functions (UDFs) were engineered using the *Camel Tools* morphological analyzer. These features included: Word Length Variance (f15), Characters per Paragraph (f38), Command Verb Count (f61), Average Sentence Length (f84), and a calculated Formality Score (f107).

1.3. Detailed Stylometric Feature Formulation

To implement a deep stylistic audit, five custom linguistic and structural features were formulated and executed via PySpark User-Defined Functions (UDFs). The breakdown of how each feature was calculated and its purpose is detailed below:

- **Word Length Variance (\$f_{15}\$):** This feature measures the structural diversity of vocabulary choices. We calculated it by finding the average word length in a text and then measuring how much individual words differ from this average. Human writers naturally mix short words with complex long words, creating higher variance, whereas AI models tend to use words of a more uniform and predictable length.
- **Characters per Paragraph (\$f_{38}\$):** This metric tracks the physical size of text blocks. We simply divided the total number of characters in the document by the total number of paragraphs. This is highly effective because AI language models typically output highly standardized, equally sized paragraphs, unlike humans who write paragraphs of varying lengths depending on the flow of ideas.
- **Command Verb Count (\$f_{61}\$):** We utilized the *Camel Tools* morphological analyzer to scan the text and count the exact number of imperative (command) verbs. Interestingly, the data analysis revealed that this feature had a value of zero across all texts. This perfectly reflects the nature of academic abstracts, which use objective, declarative language rather than

commands. The model successfully learned to ignore this feature and focus on other metrics.

- **Average Sentence Length (\$f_{84}\$):** To capture sentence structure, we counted the total number of words in a text and divided it by the number of sentences (identified by standard Arabic punctuation marks). This emerged as our most powerful feature, as it successfully caught the mechanical, consistent sentence lengths generated by AI, comparing them against the natural, fluctuating sentence lengths of human writers.
- **Formality Score (\$f_{107}\$):** This metric evaluates the academic tone of the text based on its Parts of Speech (POS). Using *Camel Tools*, we calculated the density of formal words (such as nouns, adjectives, and prepositions) and compared it against the number of narrative words (such as verbs and pronouns). A higher concentration of nouns and adjectives indicates a higher formality score, helping the model detect when an AI synthesizes an overly rigid or formal academic style.

1.4. Pipeline Optimization and Data Persistence

Processing unstructured Arabic text and running morphological disambiguation on 17,000 records is computationally expensive. To optimize the Big Data pipeline and mitigate memory limitations within the cluster, the engineered features were persisted to the disk. By writing the intermediate DataFrame to a distributed **Parquet** format, the Apache Spark lineage (DAG) was effectively broken. This strategic persistence eliminated the computational bottleneck, reducing redundant processing times and allowing for rapid, iterative model training. Finally, the TF-IDF and stylometric features were concatenated using a VectorAssembler and standardized via a StandardScaler to ensure uniform scale across all dimensions.

2. Results and Discussion

2.1. Model Performance and Evaluation After completing the distributed feature engineering process, the dataset was partitioned into a 70% training set, 15% validation set, and 15% testing set. Three machine learning architectures were trained and evaluated to identify the most robust model for detecting AI-generated Arabic academic text. The baseline performance was established using the Naïve Bayes classifier, followed by advanced testing using Random Forest and Gradient Boosted Trees (GBT).

The empirical accuracy scores obtained during the validation phase are summarized in the table below:

Classification Model	Validation Accuracy	Role in the Pipeline
Naïve Bayes	88.86%	Baseline Performance Benchmark
Gradient Boosted Trees (GBT)	89.76%	Iterative Advanced Model
Random Forest	90.78%	Selected Champion Model

To provide a visual synthesis of the model selection process, a performance comparison chart was generated. This visual representation highlights the resilience of advanced ensemble methods over standard baseline frameworks. The Random Forest classifier stood out as the undisputed champion, being the only model to break the 90% accuracy threshold, demonstrating its superior ability to handle this specific high-dimensional feature space without overfitting.




Figure 1: Model Accuracy Comparison Chart

2.2. Classification Matrix and Error Inferences To deeply investigate the classification integrity of the champion Random Forest model, a confusion matrix was computed. Evaluating the model beyond simple accuracy is essential to ensure that the detector maintains statistical equilibrium and does not suffer from a class-propensity bias (such as over-classifying texts as purely human or purely AI).




Figure 2: Random Forest Confusion Matrix Heatmap

The resulting matrix demonstrates an exceptional capability to differentiate between both target classes. The balanced true-positive and true-negative distribution confirms that the combination of lexical tokens and custom structural attributes provides stable boundary conditions for classification.

According to the comprehensive prediction ledger exported to the audit file (Full_Dataset_Predictions.xlsx), the marginal errors committed by the model were primarily linked to human abstracts written in an excessively rigid, passive, and formal manner. This human stylistic choice mimicked the standardized paragraph densities and elevated formality scores typically produced by large language models, leading to occasional false positives.

2.3. Stylometric Feature Importance and Linguistic Analysis A critical objective of this research was to understand how the engineered features contribute to the model's final decision-making process. While the collective 2,000 TF-IDF features establish the broad contextual vocabulary framework, isolating the custom stylometric features reveals profound insights into the structural habits of artificial intelligence.

The relative importance of each engineered custom feature was computed and mapped below to identify the most potent indicators of machine-generated text:




Figure 3: Stylometric Feature Importance Bar Chart

The results deliver a powerful confirmation of the original research hypothesis:

- **Average Sentence Length (f84):** Emerged as the dominant predictor, accounting for **52.4%** of the relative predictive weight among custom features. This indicates that the mathematical consistency and automated pacing of AI sentence boundaries create a distinct statistical contrast against the fluctuating sentence lengths of human writers.
- **Characters per Paragraph (f38):** Ranked second at **38.8%**, capturing the mechanical tendency of language models to generate perfectly balanced, dense block paragraphs.
- **Formality Score (f107):** Provided a valuable stylistic gradient (**7.3%**), mapping the overly formal or noun-dense structures characteristic of synthetic academic translations.
- **Word Length Variance (f15):** Contributed a measurable **1.5%** to the model's decision. While minor compared to broader syntactic features, it successfully provided a subtle signal regarding the natural dispersion and diversity of human vocabulary lengths versus the more uniform word choices of LLMs.
- **Command Verb Count (f61):** Formally recorded **0.0%** variance across the corpus. Far from being a failure, this finding proves that the system successfully recognized the domain-specific bounds of academic text, which is naturally declarative and entirely objective.

3. Real-Time Streaming Simulation

To demonstrate the deployment viability of the developed AI text detection architecture, a real-time streaming pipeline was simulated using **Spark Structured Streaming**. First, the trained champion Random Forest model and the standardization scaler were persisted to the distributed environment to enable immediate, on-the-fly inference without the need for model retraining.

A streaming DataFrame was then initialized using a structured schema to continuously monitor a dedicated landing directory for incoming JSON data streams. The pipeline utilized a micro-batch processing engine configured in append mode to process text flows seamlessly as they arrived. To evaluate the pipeline under operational conditions, an automated producer script was executed to generate and drop mock text messages into the ingestion directory at regular five-second intervals.

The streaming consumer successfully intercepted, parsed, and classified the incoming streaming data within the live console interface. This implementation confirms that the developed Big Data architecture is fully capable of conducting live, real-time academic fraud detection and automated content auditing in a production-ready environment.

```
check_stream = spark.read.json("stream_input")
check_stream.select("text", "label").show(10, truncate=70)
```

	text	label
10	... تبتذلها الجرائ...	
10	... اتبدو العلقه بين الفلسفه والعلوم الانسانيه علاقه ذات مسارب مختلفه ا...	
10	... ادفع تفاقم الانكماش الاقتصادي العالمي بالسلطات النقدية لاهم البلدان ...	
10	... اتعد ولايه البويره التي تقع جنوب شرق الجزائر العاصمه من المناطق التي...	
10	... اتسم الكتابه العلميه بالحكامه والجوده والدقه والمسؤوليه الاخلاقيه و...	

Figure 4: Streaming Console Output Screenshot

4. Conclusion and Future Work

This project successfully developed an end-to-end, high-performance Big Data pipeline tailored for the detection of AI-generated Arabic academic text. By leveraging Apache Spark's distributed computing framework, a substantial corpus of 16,895 balanced records was efficiently processed, cleaned, and analyzed, overcoming traditional computational bottlenecks through strategic Parquet storage persistence.

The integration of foundational lexical vectors via TF-IDF with targeted, custom stylometric features proved highly effective. The ensemble architecture of the **Random Forest** classifier achieved a superior validation accuracy of **90.78%**, validating the underlying hypothesis that machine-generated Arabic prose leaves a distinct and quantifiable structural footprint. The analytical deep dive established that **Average Sentence Length (f_84)** was the most dominant individual stylistic predictor, capturing 52.4% of the relative custom feature weight. Furthermore, the deployment of a functional Spark Structured Streaming pipeline demonstrated the operational readiness of the system for live academic auditing.

Future extensions of this research will focus on expanding the dataset to incorporate diverse academic disciplines, testing the pipeline against newer iterations of large language models, and exploring the integration of deep learning architectures within the distributed Spark environment to further elevate detection thresholds.